

MEHUL GOYAL

Software Engineer · BTech(Hons.) CSE - AI/ML

✉ mehulgoyal238@gmail.com — 🌐 github.com/mehulgoyal — in linkedin.com/in/mehulgoyal — 🌐 mehulgoyal.vercel.app

Education

B.Tech (Hons) CSE – AI/ML, Minor in Fintech

RV University, Bengaluru

2023 – 2027

GPA: 9.01 / 10.00

Work Experience

Machine Learning Trainee — Tata Elxsi

Jun – Sep 2025

Industry Client Project

- ▶ Engineered a customer segmentation pipeline over 50K+ records, enabling targeting of 4 distinct cohorts and reducing estimated campaign waste by **30%**.
- ▶ Built a production RAG pipeline (LangChain + local LLMs, zero API cost) cutting analyst decision-support time from hours to **under 2 minutes** per query.
- ▶ Automated 3 manual cross-team reporting steps via autonomous AI agents, eliminating recurring overhead.

Machine Learning Intern — RV University

Jun – Sep 2025

Computer Vision · Disaster Management

- ▶ Delivered end-to-end disaster CV pipeline (YOLOv11): **91%+ mAP** across 6 disaster categories, integrated into a live emergency-response dashboard.
- ▶ Cut inference latency by **38%** via optimized preprocessing and batched inference — enabling real-time edge deployment.

AI/ML Intern — Dhee Centre of AI & Data Science

Jun – Sep 2024

Model Development · Data Analysis

- ▶ Delivered 3 client-facing ML models; one reached production for student outcome prediction, actively used by the institute.
- ▶ Led a 4-person team to deliver 2 weeks ahead of schedule, earning commendation from the department head.

For more projects visit: [Projects](#)

Featured Projects

Agentic Bias Checkpoint Auditor

[GitHub](#)

Python · LangChain · Scikit-learn · Open-source LLMs

- ▶ Audits ML pipelines at 3 checkpoints (Data → Features → Model) autonomously — a gap unaddressed by existing OSS tools; adopted by 2 university research groups within 4 weeks.

Agentic Scientific Debugger

[Demo](#)

Python · LangChain · Llama · Flask

- ▶ Detects methodological flaws in ML experiments and generates evidence-backed corrections using local LLMs; reduced experiment review time by **~70%** across 15 logs.

TraceMaint — BeachHack S7, 4th / 400+ teams (Top 1%)

[GitHub](#)

Python · Constrained LLMs · Decision Trees · FastAPI

- ▶ Explainable predictive maintenance with fully auditable decision traces; cut false-positive alerts by **44%** vs. baseline LSTM in demo evaluation.

Technical Skills

Languages	Python, C++, SQL, Solidity
AI / ML	PyTorch, TensorFlow, Scikit-learn, OpenCV, LangChain
Web & Backend	Next.js, React, Flask, FastAPI, PostgreSQL, Firebase, REST APIs
DevOps	Git, Docker, GitHub Actions, CI/CD, Vercel
Other	Qiskit (Quantum), Solidity / Web3

Awards & Certifications

Awards

- ▶ Winner, Best Project — SilentSync · IEEE Pre-Conference
- ▶ 2nd Place — IEEE RAS Conference 2025
- ▶ Top-1% Finalist — BeachHack S7 2026 (TraceMaint)
- ▶ Finalist — M-Hash Hackathon 2024 (PetCompass)

Certifications

- ▶ IBM Qiskit (IUC EE)
- ▶ IBM SkillsBuild: GenAI & Big Data
- ▶ Agentic AI — Google
- ▶ Salesforce Agentblazer Champion '25